

Recurrent Neural Networks

Kate Farrahi

ECS Southampton

March 7, 2019

Batch Normalisation

- ▶ Batch normalization (BN) is a technique for improving the performance of ANNs
- ▶ It improves the stability of ANNs by adjusting and scaling the activations
- ▶ BN makes your ANN more robust to the choice of hyperparameters (larger range of parameters that will work well)
- ▶ It was introduced by Sergey Ioffe and Christian Szegedy in 2015 ¹

¹<https://arxiv.org/pdf/1502.03167.pdf>

Batch Normalisation

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_{1\dots m}\}$;
Parameters to be learned: γ, β

Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

Algorithm 1: Batch Normalizing Transform, applied to activation x over a mini-batch.

Batch Normalisation

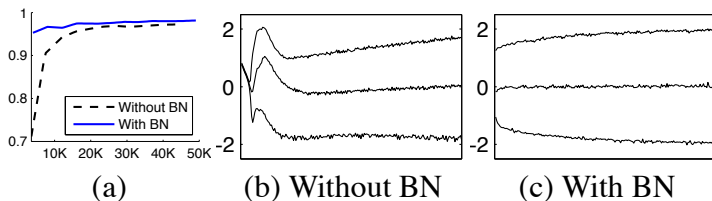


Figure 1: (a) *The test accuracy of the MNIST network trained with and without Batch Normalization, vs. the number of training steps. Batch Normalization helps the network train faster and achieve higher accuracy.* (b, c) *The evolution of input distributions to a typical sigmoid, over the course of training, shown as $\{15, 50, 85\}$ th percentiles. Batch Normalization makes the distribution more stable and reduces the internal covariate shift.*

Batch Normalisation

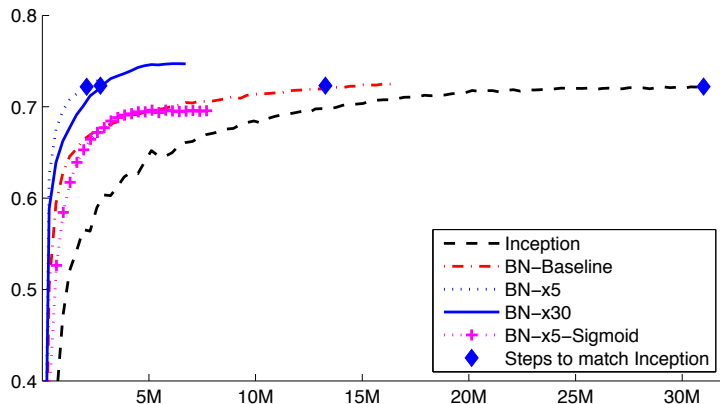


Figure 2: *Single crop validation accuracy of Inception and its batch-normalized variants, vs. the number of training steps.*

Recurrent Neural Networks - Motivation

x : Kate Farrahi and Jonathon Hare teach deep learning

y : 1 1 0 1 1 0 0 0

Recurrent Neural Networks - Motivation

$x:$	$x^{<1>}$	$x^{<2>}$...	$x^{<t>}$...	$x^{<T_x>}$
$x:$	Kate	Farrahi	...	Hare	...	learning
$y:$	$y^{<1>}$	$y^{<2>}$...	$y^{<t>}$...	$y^{<T_y>}$
$y:$	1	1	...	1	...	0

In this example, $T_x = T_y = 8$ but T_x and T_y can be different.

One Hot Encoding

How can we represent individual words?

"a"	"abbreviations"		"zoology"	"zoom"
1	0		0	0
0	1		0	1
0	0		0	0
.
.
.
0	0		0	0
0	0		1	0
0	0		0	1

2

²<https://ayearofai.com>

Why Not a Standard Feed Forward Network?

- ▶ For a task such as "Named Entity Recognition" a feed forward network would have several disadvantages

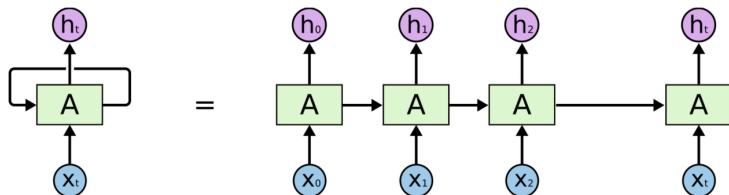
Why Not a Standard Feed Forward Network?

- ▶ For a task such as "Named Entity Recognition" a feed forward network would have several disadvantages
- ▶ The inputs and outputs may have varying lengths

Why Not a Standard Feed Forward Network?

- ▶ For a task such as "Named Entity Recognition" a feed forward network would have several disadvantages
- ▶ The inputs and outputs may have varying lengths
- ▶ The features wouldn't be shared across different positions in the network

Recurrent Neural Networks

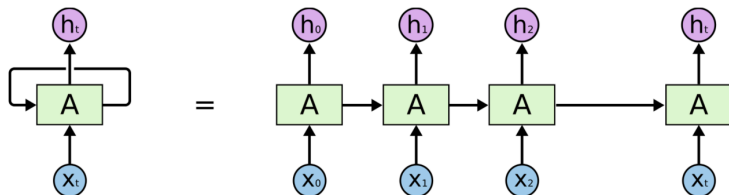


3

- ▶ RNNs are a family of ANNs for processing sequential data

³Image taken from <https://towardsdatascience.com>

Recurrent Neural Networks

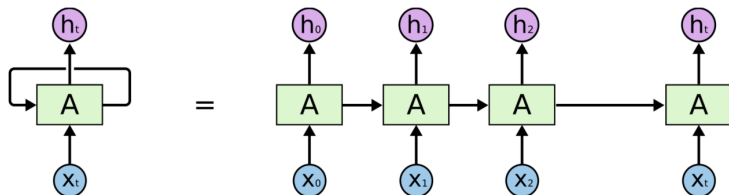


3

- ▶ RNNs are a family of ANNs for processing sequential data
- ▶ RNNs have directed cycles in their computational graphs

³Image taken from <https://towardsdatascience.com>

Recurrent Neural Networks

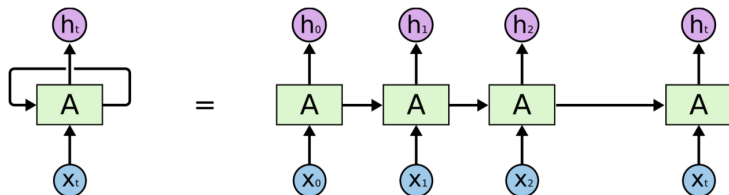


3

- ▶ RNNs are a family of ANNs for processing sequential data
- ▶ RNNs have directed cycles in their computational graphs
- ▶ They can have complicated dynamics, difficult to train

³Image taken from <https://towardsdatascience.com>

Recurrent Neural Networks



3

- ▶ RNNs are a family of ANNs for processing sequential data
- ▶ RNNs have directed cycles in their computational graphs
- ▶ They can have complicated dynamics, difficult to train
- ▶ They are more biologically realistic

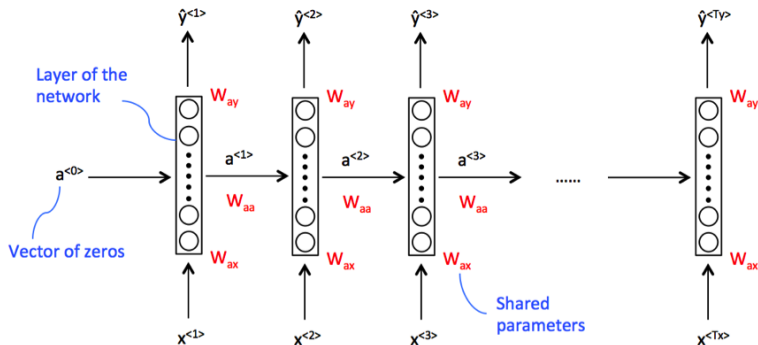
³Image taken from <https://towardsdatascience.com>

Recurrent Neural Networks

RNNs combine two properties which make them very powerful.

1. Distributed hidden state that allows them to store a lot of information about the past efficiently. This is because several different units can be active at once, allowing them to remember several things at once.
2. Non-linear dynamics that allows them to update their hidden state in complicated ways. They can however have complicated dynamics, making them difficult to train

Backpropagation Through Time (BPTT)



4

⁴Image taken from Andrew Ng

BPTT - Forward Pass

$$a^{<t>} = g(w_{aa}a^{<t-1>} + w_{ax}x^{<t>} + b_a) \quad (1)$$

BPTT - Forward Pass

$$a^{<t>} = g(w_{aa}a^{<t-1>} + w_{ax}x^{<t>} + b_a) \quad (1)$$

$$\hat{y}^{<t>} = g(w_{ya}a^{<t>} + b_y) \quad (2)$$

BPTT - Forward Pass

$$a^{<t>} = g(w_{aa}a^{<t-1>} + w_{ax}x^{<t>} + b_a) \quad (1)$$

$$\hat{y}^{<t>} = g(w_{ya}a^{<t>} + b_y) \quad (2)$$

$$\mathcal{L}^{<t>} = -y^{<t>} \log(\hat{y}^{<t>}) - (1 - y^{<t>}) \log(1 - \hat{y}^{<t>}) \quad (3)$$

BPTT - Forward Pass

$$a^{<t>} = g(w_{aa}a^{<t-1>} + w_{ax}x^{<t>} + b_a) \quad (1)$$

$$\hat{y}^{<t>} = g(w_{ya}a^{<t>} + b_y) \quad (2)$$

$$\mathcal{L}^{<t>} = -y^{<t>} \log(\hat{y}^{<t>}) - (1 - y^{<t>}) \log(1 - \hat{y}^{<t>}) \quad (3)$$

$$\mathcal{L} = \sum_{t=1}^{T_y} \mathcal{L}^{<t>} \quad (4)$$

BPTT - Backwards Pass

$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{ya}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial w_{ya}} \quad (5)$$

BPTT - Backwards Pass

$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{ya}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial w_{ya}} \quad (5)$$

$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{aa}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial \hat{a}^{<3>}} \frac{\partial \hat{a}^{<3>}}{\partial w_{aa}} \quad (6)$$

$$(7)$$

BPTT - Backwards Pass

$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{ya}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial w_{ya}} \quad (5)$$

$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{aa}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial a^{<3>}} \frac{\partial \hat{a}^{<3>}}{\partial w_{aa}} \quad (6)$$


$$(7)$$

$$\text{Recall } a^{<3>} = g(w_{aa}a^{<2>} + w_{ax}x^{<3>} + b_a)$$


$$\frac{\partial \mathcal{L}^{<3>}}{\partial w_{aa}} = \frac{\partial \mathcal{L}^{<3>}}{\partial \hat{y}^{<3>}} \frac{\partial \hat{y}^{<3>}}{\partial a^{<3>}} \frac{\partial a^{<3>}}{\partial a^{<2>}} \frac{\partial a^{<2>}}{\partial a^{<1>}} \frac{\partial a^{<1>}}{\partial w_{aa}} \quad (8)$$

Recurrent Neural Networks


one to one




one to many




many to one



many to many



many to many



5

⁵<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>